

Video Streaming Services EDA

February 29, 2024

1 Video Streaming Service Platforms - EDA

1.0.1 SUMMARY:

We were tasked with coming up with 4 interesting questions to answer using the data sets provided. Our group chose to examine the 4 largest video streaming platforms: Netflix, Hulu, Disney+, and Amazon Prime Video.

The data provided was simply a list of each platform's title catalog, including information on rating, duration, director, and more. That being said, many of these data fields needed significant cleaning, as you will see below, thus limiting our analysis to type (TV Show/Movie), title, date_added, release_year, rating, duration, and main genre. The data sets contain mostly qualitative data, as performance metrics such as total plays, total positive ratings, new user registration data were not included. This fact made us take significant time in determining who our stakeholder for this analysis could be.

Overall, we settled on the following stakeholder profile: Producer who is trying figure out which streaming platform to sell their new R-rated movie to, as well as, what content-type and genre should future content focus on.

We attempt to answer the following questions for our defined stakeholder: 1) Which streaming service has the highest volume of "new" content? 2) Which streaming service is the most likely to buy an R-rated movie? 3) Which content type is more likely to be bought? 4) Which genres should future content offerings focus on in order to increase the likelihood of being purchased by a streaming platform?

1.0.2 CONCLUSIONS:

Based on our extensive EDA below, we can attempt to answer the four questions posed: 1) Which streaming service has the highest volume of "new" content? - Based on our analysis, we can look at the distribution of title counts by release year. The platform with the highest proportion of titles from recent years compared to their full library should give us an idea which platform focuses their efforts on new content the most. In this case, Amazon Prime Video has the highest proportion of titles with 2021 release year in their library, so we can conclude that they put high value on having the newest content. 2) Which streaming service is the most likely to buy an R-rated movie? - Based on our analysis, we can look at the distribution of title counts by rating. Similar to the above question, it's a matter of proportion. Whichever platform has the highest proportion of R-rated movies in their library should signify a willingness to purchase R-rated movies. In this case, Amazon Prime Video has the highest proportion of titles with an R rating, thus showing a high willingness to purchase R-rated movies. An important caveat here is that Amazon Prime Video

has the largest library of content, showing a willingness to buy everything. 3) Which content type is more likely to be bought? - This question requires us to drill down a bit further by looking at the proportion of movies to tv shows across all streaming platforms. We can say that movies are more likely to be bought by Netflix, Disney+, and Amazon Prime Video than by Hulu. However, Hulu is the only platform that seems to balance it's offerings of movies and tv shows, showing a greater willingness to purchase tv shows than the other platforms. 4) Which genres should future content offerings focus on in order to increase the likelihood of being purchased by a streaming platform? - When looking at the distribution of genres by title count, we can get an idea of which genres are the most popular. Drama is the genre most popular on Netflix and Amazon Prime Video, while Action/Adventure take the top spot on Hulu and Disney+. To answer this question completely, we would have to answer it on a case-by-case basis for each title.

2 Import Statements

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sqlalchemy import create_engine
from sqlalchemy.sql import text
%matplotlib inline
```

3 Database Read In

First, we must read in the 4 databases we will be working with. 3 are in Excel format and 1 is in PostgreSQL.

3.0.1 Netflix DataFrame

```
[15]: netflix_df = pd.read_csv('./data/netflix_titles.csv')
```

3.0.2 Hulu DataFrame

```
[3]: hulu_df = pd.read_csv('./data/hulu_titles.csv')
```

3.0.3 Disney+ DataFrame

```
[4]: disneyplus_df = pd.read_csv('./data/disney_plus_titles.csv')
```

3.0.4 Amazon Prime Video DataFrame (FROM SQL)

```
[4]: engine = create_engine('postgresql+psycopg2://postgres:hellosql@localhost/
↳amazonprimestreaming')
```

```
[5]: amazon_sql = "SELECT * FROM amazon_prime_titles"
amazon_df = pd.read_sql(amazon_sql, engine)
```

4 Data Cleaning

We follow the same basic steps for data cleaning each individual data set. The steps we took are as follows: 1) Drop columns that we are uninterested in, such as director, cast, description - Because we cannot guarantee that specific values are correct, such as the spelling of a director's or cast member's name being consistently correct, analysis on those columns is futile.

2) Drop NA values responsibly

- We must pay attention to the number of records lost to a `.dropna()` operation. Dropping too many will ruin the validity of the data set.

3) Data Engineering and calculated column creation

- Due to multiple genres being listed per title in the `listed_in` column, we had to split them out into their own columns based on a delimiter. We ended up taking the first genre listed as the main genre and removed any secondary genres.

4.0.1 Netflix Data Cleaning

```
[16]: clean_netflix_df = netflix_df.copy()
```

```
[17]: clean_netflix_df = clean_netflix_df.drop(['director', 'cast', 'description'],  
      ↪axis=1)
```

```
[18]: clean_netflix_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 8807 entries, 0 to 8806  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   show_id         8807 non-null   object  
1   type            8807 non-null   object  
2   title           8807 non-null   object  
3   country         7976 non-null   object  
4   date_added      8797 non-null   object  
5   release_year    8807 non-null   int64  
6   rating          8803 non-null   object  
7   duration        8804 non-null   object  
8   listed_in       8807 non-null   object  
dtypes: int64(1), object(8)  
memory usage: 619.4+ KB
```

```
[19]: new_netflix_df = clean_netflix_df['listed_in'].str.split(',', expand=True)
```

```
[20]: frames = [clean_netflix_df, new_netflix_df]  
final_netflix_df = pd.concat(frames, axis=1, join="inner")
```

```
[21]: new_df = final_netflix_df.drop('listed_in', axis=1)
```

```
[22]: fin_netflix_df = new_df.rename(columns = {0: 'genre_1', 1: 'genre_2', 2: 'genre_3'})

[23]: finished_netflix_df = fin_netflix_df.drop(['genre_2', 'genre_3'], axis=1)

[15]: finished_netflix_df.to_csv('./GitHub/GC_DA_Capstone/final_data/netflix_df.csv', index = False)
```

4.0.2 Hulu Data Cleaning

```
[16]: clean_hulu_df = hulu_df.copy()

[17]: cleaner_hulu_df = clean_hulu_df.drop(['director', 'cast', 'description'], axis=1)

[18]: cleaner_hulu_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3073 entries, 0 to 3072
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   show_id         3073 non-null   object
1   type            3073 non-null   object
2   title           3073 non-null   object
3   country         1620 non-null   object
4   date_added      3045 non-null   object
5   release_year    3073 non-null   int64
6   rating          2553 non-null   object
7   duration        2594 non-null   object
8   listed_in       3073 non-null   object
dtypes: int64(1), object(8)
memory usage: 216.2+ KB

[19]: new_hulu_df = cleaner_hulu_df['listed_in'].str.split(',', expand=True)

[20]: frames = [cleaner_hulu_df, new_hulu_df]
final_hulu_df = pd.concat(frames, axis=1, join="inner")

[21]: fin_hulu_df = final_hulu_df.rename(columns = {0: 'genre_1', 1: 'genre_2', 2: 'genre_3'})

[22]: finished_hulu_df = fin_hulu_df.drop(['listed_in', 'genre_2', 'genre_3'], axis=1)

[23]: finished_hulu_df.to_csv('./GitHub/GC_DA_Capstone/final_data/hulu_df.csv', index = False)
```

4.0.3 Disney+ Data Cleaning

```
[24]: clean_disney_df = disneyplus_df.copy()
```

```
[25]: cleaner_disney_df = clean_disney_df.drop(['director', 'cast', 'description'],  
      ↪axis=1)
```

```
[26]: cleaner_disney_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1450 entries, 0 to 1449  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   show_id         1450 non-null   object  
1   type            1450 non-null   object  
2   title           1450 non-null   object  
3   country         1231 non-null   object  
4   date_added      1447 non-null   object  
5   release_year    1450 non-null   int64  
6   rating          1447 non-null   object  
7   duration        1450 non-null   object  
8   listed_in       1450 non-null   object  
dtypes: int64(1), object(8)  
memory usage: 102.1+ KB
```

```
[27]: cleaned_disney_df = cleaner_disney_df.dropna()
```

```
[28]: cleaned_disney_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 1228 entries, 2 to 1449  
Data columns (total 9 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   show_id         1228 non-null   object  
1   type            1228 non-null   object  
2   title           1228 non-null   object  
3   country         1228 non-null   object  
4   date_added      1228 non-null   object  
5   release_year    1228 non-null   int64  
6   rating          1228 non-null   object  
7   duration        1228 non-null   object  
8   listed_in       1228 non-null   object  
dtypes: int64(1), object(8)  
memory usage: 95.9+ KB
```

```
[29]: new_disney_df = cleaned_disney_df['listed_in'].str.split(',', expand=True)
```

```
[30]: frames = [cleaned_disney_df, new_disney_df]
      final_disney_df = pd.concat(frames, axis=1, join="inner")

[31]: fin_disney_df = final_disney_df.rename(columns = {0: 'genre_1', 1: 'genre_2', 2:
      ↪ 'genre_3'})

[32]: finished_disney_df = fin_disney_df.drop(['listed_in', 'genre_2', 'genre_3'],
      ↪ axis=1)

[33]: finished_disney_df.to_csv('./GitHub/GC_DA_Capstone/final_data/disney_df.csv',
      ↪ index = False)
```

4.0.4 Amazon Prime Video Data Cleaning

```
[6]: clean_amazon_df = amazon_df.copy()

[7]: cleaner_amazon_df = clean_amazon_df.drop(['director', 'cast', 'description'],
      ↪ axis=1)

[8]: cleaner_amazon_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9668 entries, 0 to 9667
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   show_id         9668 non-null   object
 1   type            9668 non-null   object
 2   title           9668 non-null   object
 3   country         672 non-null    object
 4   date_added      155 non-null    object
 5   release_year    9668 non-null   int64
 6   rating          9331 non-null   object
 7   duration        9668 non-null   object
 8   listed_in       9668 non-null   object
dtypes: int64(1), object(8)
memory usage: 679.9+ KB
```

```
[9]: new_amazon_df = cleaner_amazon_df['listed_in'].str.split(',', expand=True)

[10]: frames = [cleaner_amazon_df, new_amazon_df]
      final_amazon_df = pd.concat(frames, axis=1, join="inner")

[11]: fin_amazon_df = final_amazon_df.rename(columns = {0: 'genre_1', 1: 'genre_2', 2:
      ↪ 'genre_3', 3: 'genre_4', 4: 'genre_5'})

[12]: finished_amazon_df = fin_amazon_df.drop(['listed_in', 'genre_2', 'genre_3',
      ↪ 'genre_4', 'genre_5'], axis=1)
```


```
[41]: finished_amazon_df.to_csv('./GitHub/GC_DA_Capstone/final_data/amazon_df.csv',  
    ↪ index = False)
```

5 Netflix EDA

The figure below shows the distribution of titles by release year. It attempts to answer the question: Which streaming service has the highest volume of “new” content?

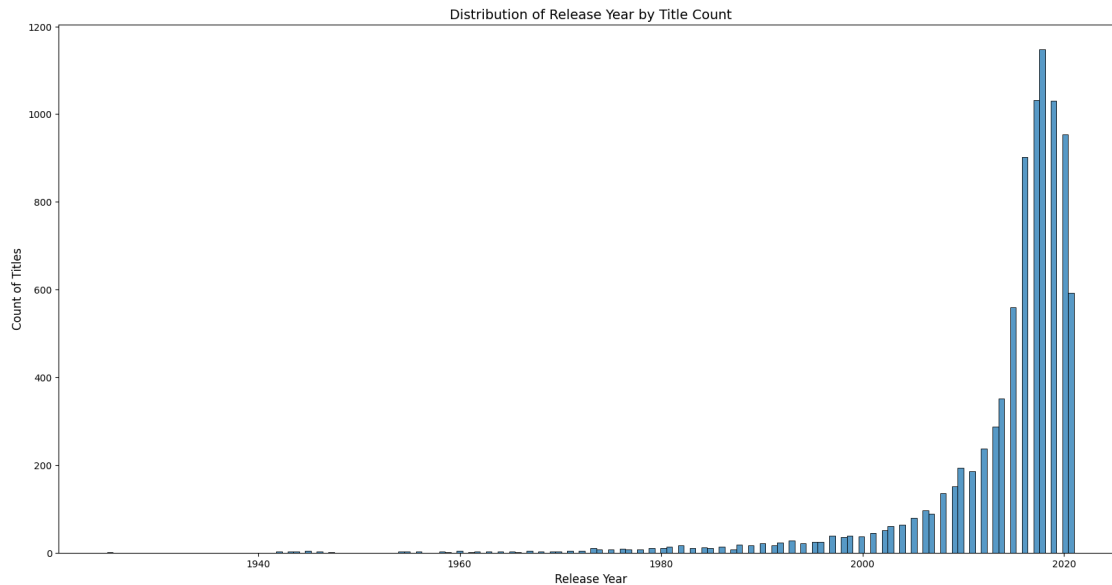
```
[42]: plt.figure(figsize = (20.0, 10.0))  
    ax = sns.histplot(finished_netflix_df, x = 'release_year', bins= 10)  
  
    plt.xlabel("Release Year (each bar equals 10 years)", fontsize = 12)  
    plt.ylabel("Count of Titles", fontsize = 12)  
    plt.title("Total Number of Titles per Decade", fontsize = 14)
```

```
[42]: Text(0.5, 1.0, 'Total Number of Titles per Decade')
```



```
[43]: plt.figure(figsize = (20.0, 10.0))  
    sns.histplot(finished_netflix_df, x = 'release_year')  
  
    plt.xlabel("Release Year", fontsize = 12)  
    plt.ylabel("Count of Titles", fontsize = 12)  
    plt.title("Distribution of Release Year by Title Count", fontsize = 14)
```

```
[43]: Text(0.5, 1.0, 'Distribution of Release Year by Title Count')
```




The figure below shows the distribution of titles by rating. It attempts to answer the question: Which streaming service is most likely to purchase an R-rated movie?

```
[44]: plt.figure(figsize=(20.0,10.0))
sns.histplot(finished_netflix_df, x = 'rating')

plt.xlabel("Title Ratings", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Rating", fontsize = 14)
```

```
[44]: Text(0.5, 1.0, 'Total Number of Titles per Rating')
```




```
[24]: finished_netflix_df['rating'].value_counts()
```

```
[24]: rating
      TV-MA      3207
      TV-14      2160
      TV-PG       863
      R          799
      PG-13       490
      TV-Y7       334
      TV-Y        307
      PG          287
      TV-G        220
      NR          80
      G           41
      TV-Y7-FV     6
      NC-17        3
      UR           3
      74 min       1
      84 min       1
      66 min       1
      Name: count, dtype: int64
```

```
[45]: genre_date_breakdown = finished_netflix_df.groupby('release_year')['genre_1'].
      ↪value_counts().reset_index()
      gd_breakdown_df = pd.DataFrame(genre_date_breakdown)
```


```
[46]: netflix_only_df = gd_breakdown_df[gd_breakdown_df['release_year'] >= 2007]
```

The figure below shows the genre breakdown per year for all titles in the library. It attempts to give us information that will help answer the question: Which genres should future content offerings focus on in order to increase the likelihood of being purchased by a streaming platform?

```
[47]: plt.figure(figsize = (20.0, 10.0))
      sns.barplot(netflix_only_df, x = 'release_year', y = 'count', hue = 'genre_1')

      plt.xlabel("Release Year", fontsize = 12)
      plt.ylabel("Count of Titles", fontsize = 12)
      plt.title("Total Number of Titles per Genre for each Release Year", fontsize = 14)
      ↪14)
```

```
[47]: Text(0.5, 1.0, 'Total Number of Titles per Genre for each Release Year')
```



```
[48]: genre_counts = finished_netflix_df['genre_1'].value_counts()
genre_df = pd.DataFrame(genre_counts)
display(genre_df)
```

genre_1	count
Dramas	1600
Comedies	1210
Action & Adventure	859
Documentaries	829
International TV Shows	774
Children & Family Movies	605
Crime TV Shows	399
Kids' TV	388
Stand-Up Comedy	334
Horror Movies	275
British TV Shows	253
Docuseries	221
Anime Series	176
International Movies	128
TV Comedies	120
Reality TV	120
Classic Movies	80
TV Dramas	67
Thrillers	65
Movies	57
TV Action & Adventure	40
Stand-Up Comedy & Talk Shows	34

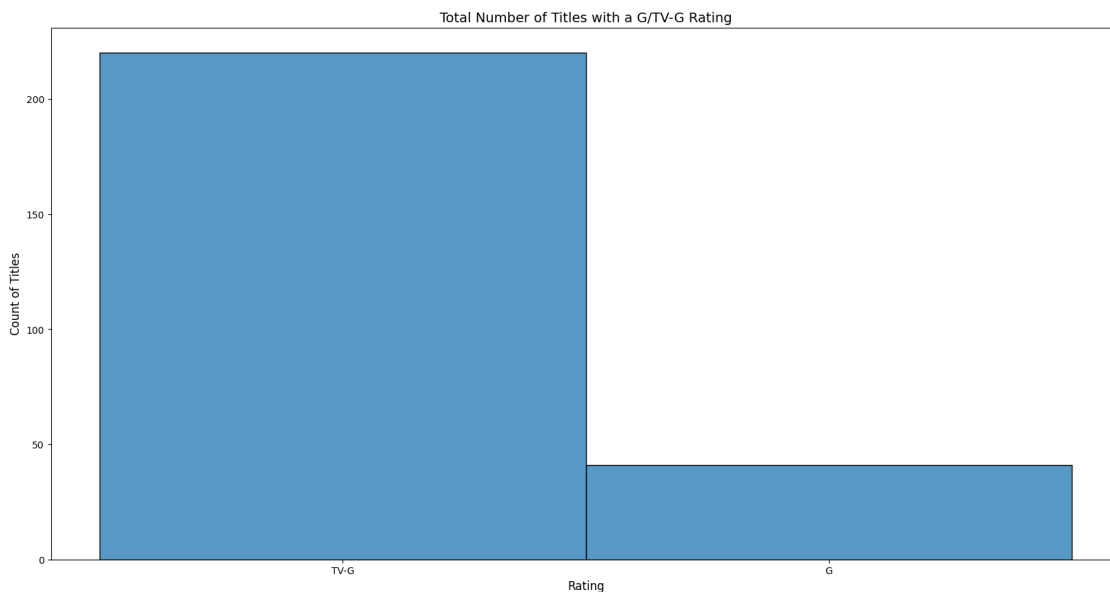
Romantic TV Shows	32
Classic & Cult TV	22
Anime Features	21
Independent Movies	20
Music & Musicals	18
TV Shows	16
Sci-Fi & Fantasy	13
Cult Movies	12
TV Horror	11
Romantic Movies	3
Spanish-Language TV Shows	2
LGBTQ Movies	1
TV Sci-Fi & Fantasy	1
Sports Movies	1

```
[49]: g_tv_g_netflix_only_df = finished_netflix_df[(finished_netflix_df['rating'] == 'TV-G') | (finished_netflix_df['rating'] == 'G')]
```

```
[50]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(g_tv_g_netflix_only_df, x = 'rating')

plt.xlabel("Rating", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles with a G/TV-G Rating", fontsize = 14)
```

```
[50]: Text(0.5, 1.0, 'Total Number of Titles with a G/TV-G Rating')
```




```
[51]: g_tv_g_netflix_only_df['rating'].value_counts()
```

```
[51]: rating
      TV-G    220
      G       41
      Name: count, dtype: int64
```

The figure below shows the proportion of movies and tv shows currently on the platform. It gives us some information we can use to answer the question: Which content type is more likely to be bought?

```
[52]: plt.figure(figsize = (20.0, 10.0))
      sns.histplot(finished_netflix_df, x='type')
```

```
[52]: <Axes: xlabel='type', ylabel='Count'>
```



```
[53]: finished_netflix_df['type'].value_counts()
```

```
[53]: type
      Movie    6131
      TV Show  2676
      Name: count, dtype: int64
```

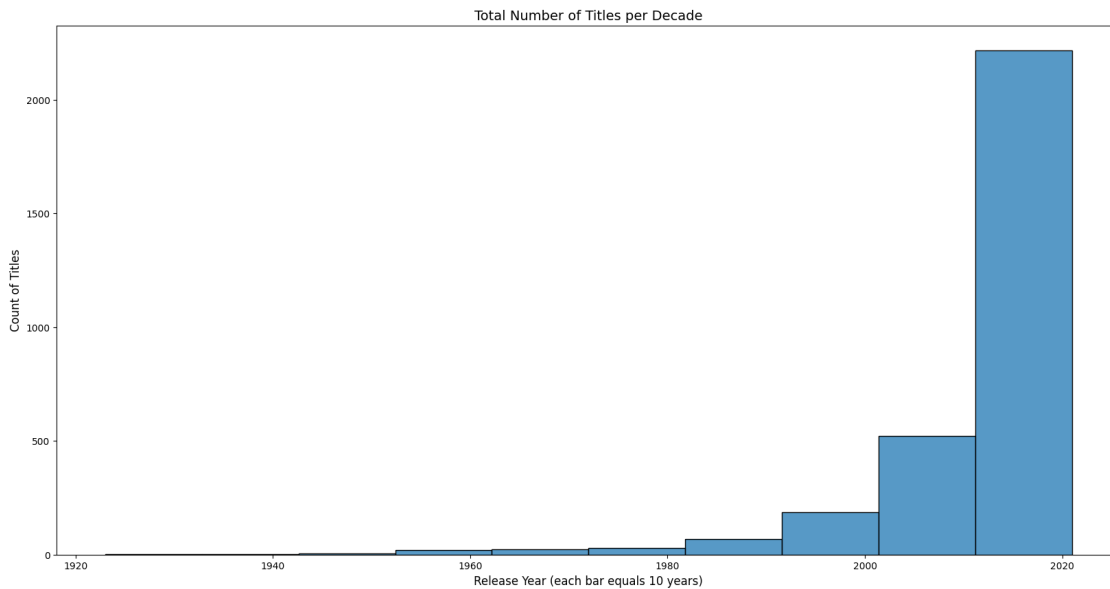
6 Hulu EDA

The EDA for each data set follows the steps taken in the Netflix EDA section. All tables and charts show the same concepts, just with different data. We can then compare the results shown in the figure to one another to get an idea if there are key differences between streaming platforms in regards to the questions asked.

```
[54]: plt.figure(figsize = (20.0, 10.0))
ax = sns.histplot(finished_hulu_df, x = 'release_year', bins= 10)

plt.xlabel("Release Year (each bar equals 10 years)", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Decade", fontsize = 14)
```

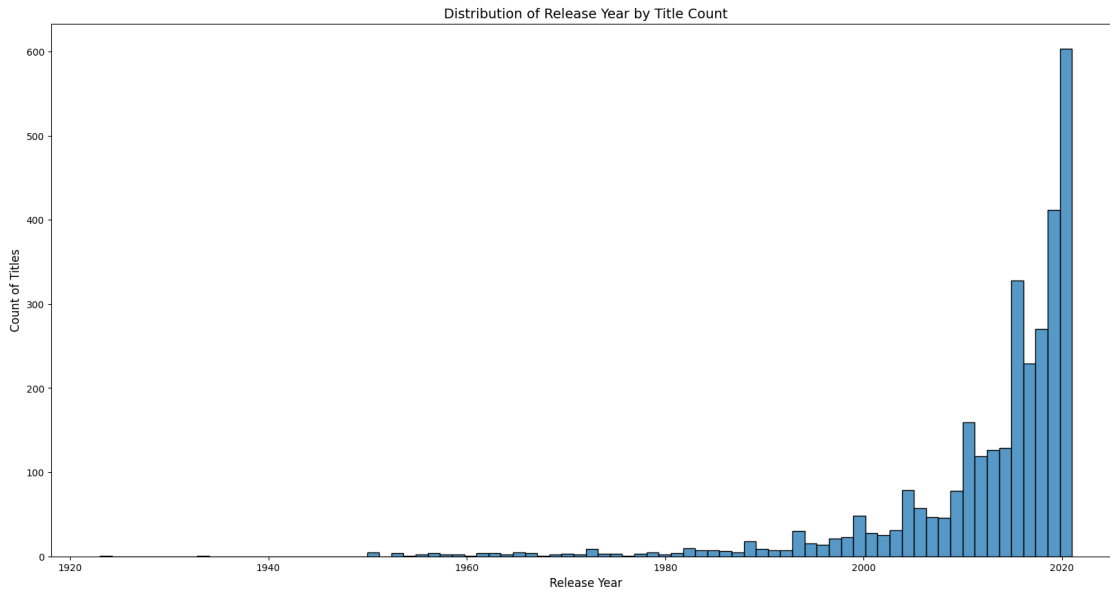
```
[54]: Text(0.5, 1.0, 'Total Number of Titles per Decade')
```



```
[55]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(finished_hulu_df, x = 'release_year')

plt.xlabel("Release Year", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Distribution of Release Year by Title Count", fontsize = 14)
```

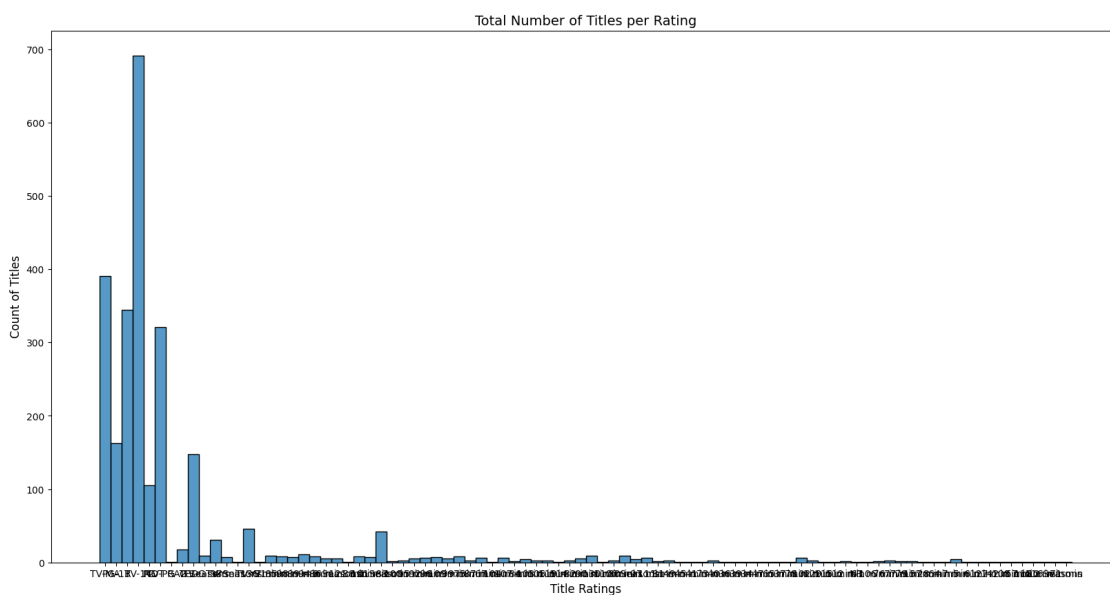
```
[55]: Text(0.5, 1.0, 'Distribution of Release Year by Title Count')
```



```
[56]: plt.figure(figsize=(20.0,10.0))
sns.histplot(finished_hulu_df, x = 'rating')

plt.xlabel("Title Ratings", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Rating", fontsize = 14)
```

```
[56]: Text(0.5, 1.0, 'Total Number of Titles per Rating')
```



```
[57]: finished_hulu_df['rating'].value_counts()
```

```
[57]: rating
      TV-14      691
      TV-MA      391
      R          345
      TV-PG      321
      PG-13      163
      ...
      34 min      1
      47 min      1
      65 min      1
      37 min      1
      71 min      1
      Name: count, Length: 88, dtype: int64
```


```
[58]: genre_date_breakdown = finished_hulu_df.groupby('release_year')['genre_1'].
      ↪value_counts().reset_index()
      gd_breakdown_df = pd.DataFrame(genre_date_breakdown)
```

```
[59]: hulu_only_df = gd_breakdown_df[gd_breakdown_df['release_year'] >= 2007]
```

```
[60]: plt.figure(figsize = (20.0, 10.0))
      sns.barplot(hulu_only_df, x = 'release_year', y = 'count', hue = 'genre_1')

      plt.xlabel("Release Year", fontsize = 12)
      plt.ylabel("Count of Titles", fontsize = 12)
      plt.title("Total Number of Titles per Genre for each Release Year", fontsize = 14)
      ↪14)
```

```
[60]: Text(0.5, 1.0, 'Total Number of Titles per Genre for each Release Year')
```



```
[61]: genre_counts = finished_hulu_df['genre_1'].value_counts()
genre_df = pd.DataFrame(genre_counts)
display(genre_df)
```

genre_1	count
Action	555
Comedy	468
Documentaries	433
Drama	415
Crime	162
Horror	149
Anime	131
Black Stories	105
Adventure	81
Cooking & Food	80
Reality	74
Sports	74
News	65
Family	49
Lifestyle & Culture	43
Kids	41
Adult Animation	29
Classics	23
Game Shows	15
Cartoons	14
Thriller	13
Latino	12


Science Fiction	10
International	7
Romance	6
Science & Technology	6
Music	6
History	3
Mystery	2
Health & Wellness	2

```
[62]: g_tv_g_hulu_only_df = finished_hulu_df[(finished_hulu_df['rating'] == 'G') |
↳ (finished_hulu_df['rating'] == 'TV-G')]
```

```
[63]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(g_tv_g_hulu_only_df, x = 'rating')

plt.xlabel("Rating", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles with a G/TV-G Rating", fontsize = 14)
```

```
[63]: Text(0.5, 1.0, 'Total Number of Titles with a G/TV-G Rating')
```




```
[64]: g_tv_g_hulu_only_df['rating'].value_counts()
```

```
[64]: rating
TV-G    148
G        18
Name: count, dtype: int64
```

```
[65]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(finished_hulu_df, x='type')
```

```
[65]: <Axes: xlabel='type', ylabel='Count'>
```




7 Disney+ EDA

```
[66]: plt.figure(figsize = (20.0, 10.0))
ax = sns.histplot(finished_disney_df, x = 'release_year', bins= 10)

plt.xlabel("Release Year (each bar equals 10 years)", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Decade", fontsize = 14)
```

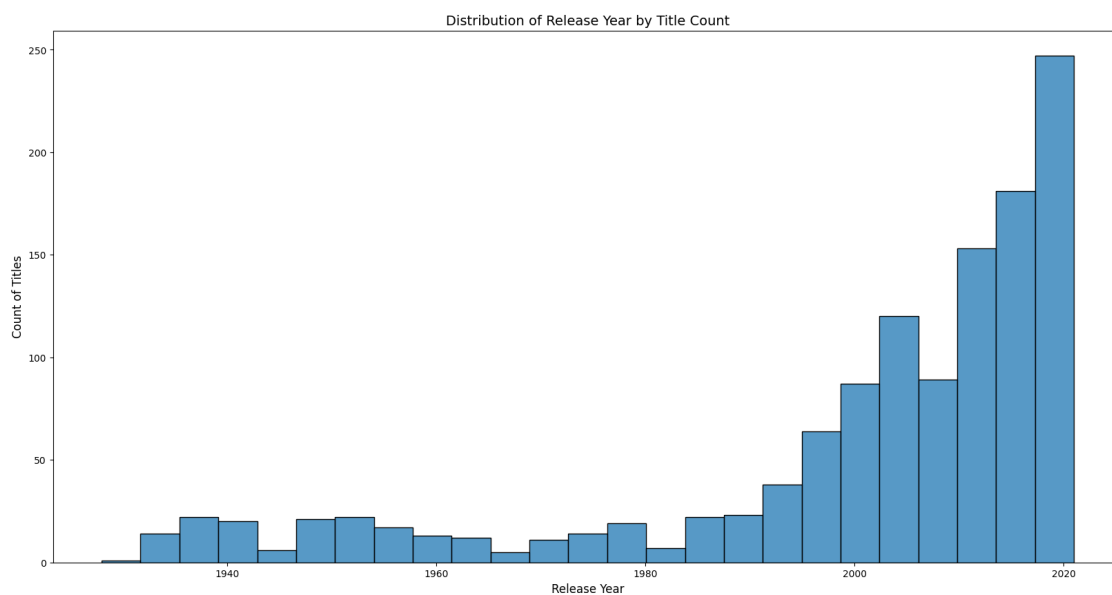
```
[66]: Text(0.5, 1.0, 'Total Number of Titles per Decade')
```



```
[67]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(finished_disney_df, x = 'release_year')

plt.xlabel("Release Year", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Distribution of Release Year by Title Count", fontsize = 14)
```

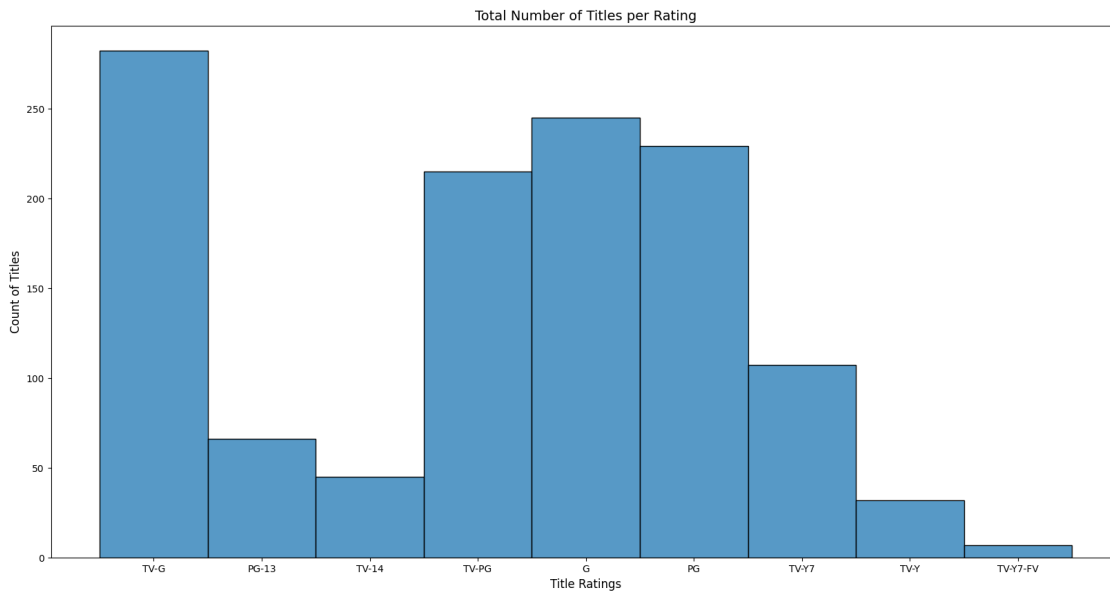
```
[67]: Text(0.5, 1.0, 'Distribution of Release Year by Title Count')
```



```
[68]: plt.figure(figsize=(20.0,10.0))
sns.histplot(finished_disney_df, x = 'rating')

plt.xlabel("Title Ratings", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Rating", fontsize = 14)
```

```
[68]: Text(0.5, 1.0, 'Total Number of Titles per Rating')
```



```
[69]: finished_disney_df['rating'].value_counts()
```

```
[69]: rating
TV-G      282
G         245
PG        229
TV-PG     215
TV-Y7     107
PG-13      66
TV-14      45
TV-Y       32
TV-Y7-FV    7
Name: count, dtype: int64
```


```
[70]: genre_date_breakdown = finished_disney_df.groupby('release_year')['genre_1'].
      ↪value_counts().reset_index()
gd_breakdown_df = pd.DataFrame(genre_date_breakdown)
```

```
[71]: disney_only_df = gd_breakdown_df[gd_breakdown_df['release_year'] >= 2007]
```

```
[72]: plt.figure(figsize = (20.0, 10.0))
sns.barplot(disney_only_df, x = 'release_year', y = 'count', hue = 'genre_1')

plt.xlabel("Release Year", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Genre for each Release Year", fontsize = 14)
```

```
[72]: Text(0.5, 1.0, 'Total Number of Titles per Genre for each Release Year')
```



```
[73]: genre_counts = finished_disney_df['genre_1'].value_counts()
genre_df = pd.DataFrame(genre_counts)
display(genre_df)
```

genre_1	count
Action-Adventure	408
Animation	280
Comedy	180
Animals & Nature	113
Coming of Age	53
Documentary	51
Biographical	31
Docuseries	25
Drama	24
Buddy	19
Family	13
Anthology	7

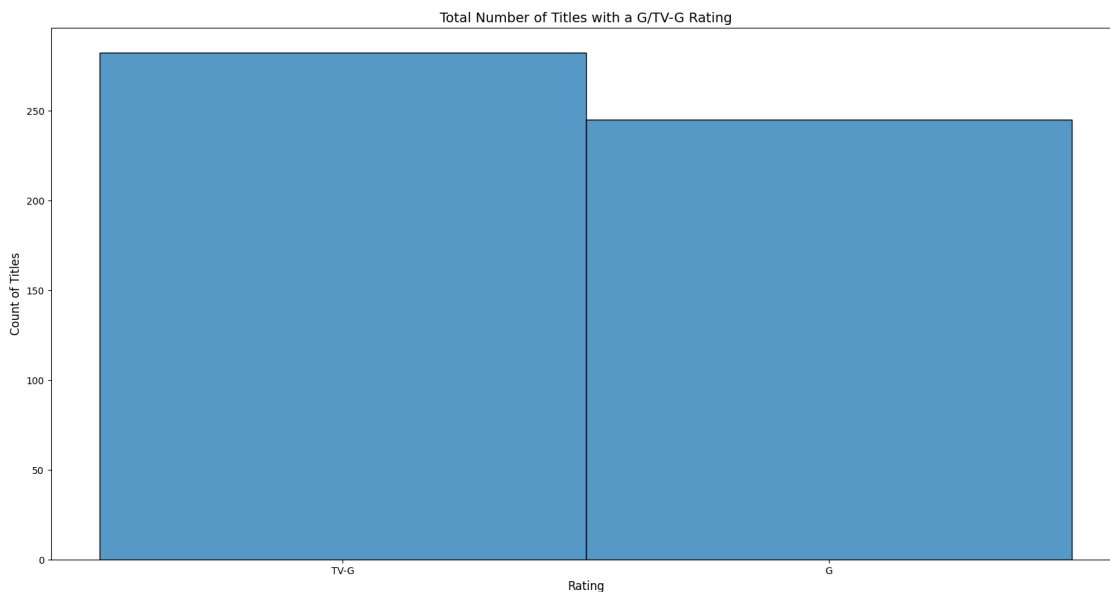
Concert Film	5
Fantasy	5
Crime	3
Historical	2
Reality	2
Movies	1
Dance	1
Musical	1
Sports	1
Music	1
Kids	1
Series	1

```
[74]: g_tv_g_disney_only_df = finished_disney_df[(finished_disney_df['rating'] == 'G') | (finished_disney_df['rating'] == 'TV-G')]
```

```
[75]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(g_tv_g_disney_only_df, x = 'rating')

plt.xlabel("Rating", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles with a G/TV-G Rating", fontsize = 14)
```

```
[75]: Text(0.5, 1.0, 'Total Number of Titles with a G/TV-G Rating')
```




```
[76]: g_tv_g_disney_only_df['rating'].value_counts()
```

```
[76]: rating
      TV-G    282
      G      245
      Name: count, dtype: int64
```

```
[77]: plt.figure(figsize = (20.0, 10.0))
      sns.histplot(finished_disney_df, x='type')
```

```
[77]: <Axes: xlabel='type', ylabel='Count'>
```




8 Amazon Prime Video EDA

```
[78]: plt.figure(figsize = (20.0, 10.0))
      ax = sns.histplot(finished_amazon_df, x = 'release_year', bins= 10)

      plt.xlabel("Release Year (each bar equals 10 years)", fontsize = 12)
      plt.ylabel("Count of Titles", fontsize = 12)
      plt.title("Total Number of Titles per Decade", fontsize = 14)
```

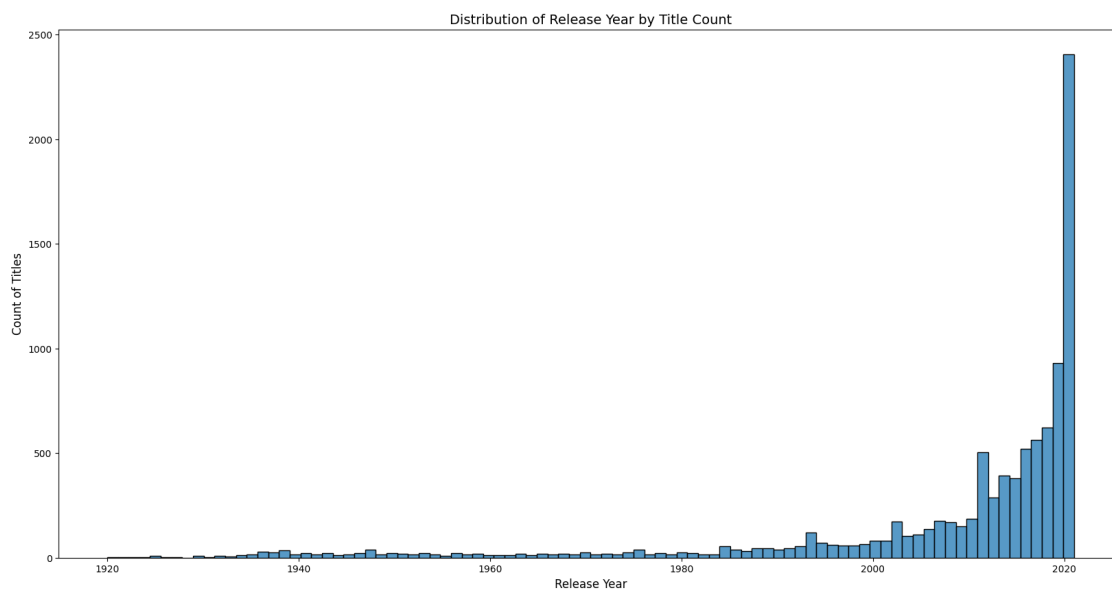
```
[78]: Text(0.5, 1.0, 'Total Number of Titles per Decade')
```



```
[79]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(finished_amazon_df, x = 'release_year')

plt.xlabel("Release Year", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Distribution of Release Year by Title Count", fontsize = 14)
```

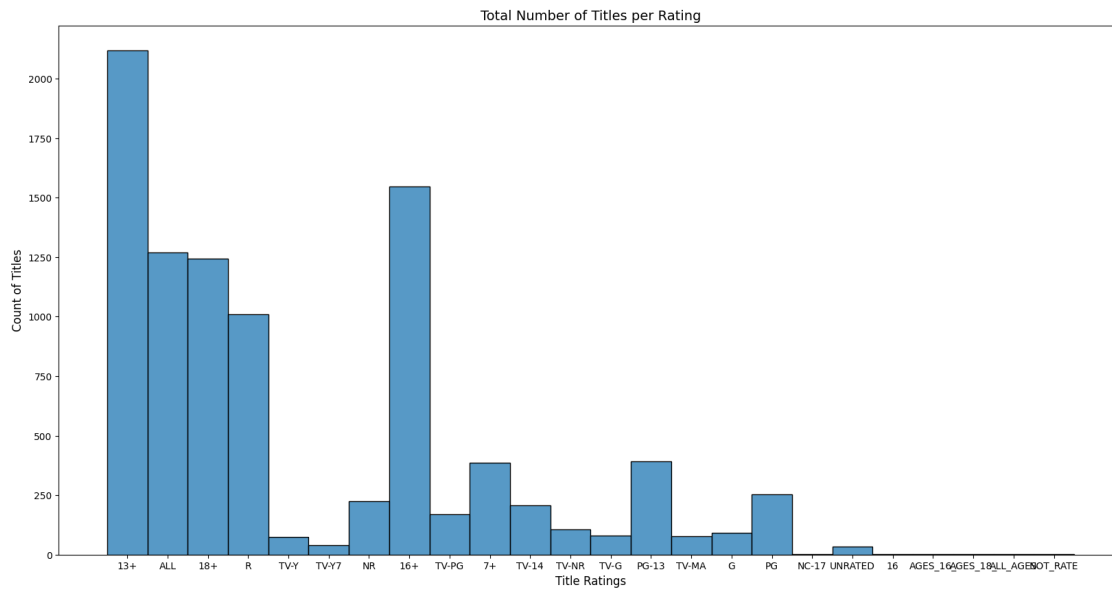
```
[79]: Text(0.5, 1.0, 'Distribution of Release Year by Title Count')
```




```
[80]: plt.figure(figsize=(20.0,10.0))
sns.histplot(finished_amazon_df, x = 'rating')

plt.xlabel("Title Ratings", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles per Rating", fontsize = 14)
```

```
[80]: Text(0.5, 1.0, 'Total Number of Titles per Rating')
```



```
[81]: finished_amazon_df['rating'].value_counts()
```

```
[81]: rating
13+      2117
16+      1547
ALL       1268
18+      1243
R         1010
PG-13     393
7+        385
PG         253
NR         223
TV-14     208
TV-PG     169
TV-NR     105
G          93
TV-G       81
TV-MA      77
TV-Y       74
```

```

TV-Y7          39
UNRATED        33
NC-17          3
AGES_18_       3
NOT_RATE       3
AGES_16_       2
16             1
ALL_AGES       1
Name: count, dtype: int64

```

```

[82]: genre_date_breakdown = finished_amazon_df.groupby('release_year')['genre_1'].
      ↪ value_counts().reset_index()
gd_breakdown_df = pd.DataFrame(genre_date_breakdown)

```

```

[83]: amazon_only_df = gd_breakdown_df[gd_breakdown_df['release_year'] >= 2007]

```

```

[84]: plt.figure(figsize = (20.0, 10.0))
      sns.barplot(amazon_only_df, x = 'release_year', y = 'count', hue = 'genre_1')


      plt.xlabel("Release Year", fontsize = 12)
      plt.ylabel("Count of Titles", fontsize = 12)
      plt.title("Total Number of Titles per Genre for each Release Year", fontsize = 14)

```

```

[84]: Text(0.5, 1.0, 'Total Number of Titles per Genre for each Release Year')

```



```
[85]: genre_counts = finished_amazon_df['genre_1'].value_counts()
genre_df = pd.DataFrame(genre_counts)
display(genre_df)
```


genre_1	count
Drama	2216
Action	1657
Comedy	1475
Documentary	913
Horror	535
Animation	498
Arts	457
Kids	373
TV Shows	263
Suspense	194
Special Interest	188
Arthouse	132
Romance	126
Music Videos and Concerts	103
Western	102
Science Fiction	85
Fitness	83
Adventure	71
International	47
Anime	44
Unscripted	29
Sports	19
Fantasy	18
Faith and Spirituality	13
LGBTQ	13
Military and War	5
Young Adult Audience	3
Talk Show and Variety	3
Historical	3

```
[86]: g_tv_g_amazon_only_df = finished_disney_df[(finished_disney_df['rating'] == 'G') | (finished_disney_df['rating'] == 'TV-G')]
```

```
[87]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(g_tv_g_amazon_only_df, x = 'rating')

plt.xlabel("Rating", fontsize = 12)
plt.ylabel("Count of Titles", fontsize = 12)
plt.title("Total Number of Titles with a G/TV-G Rating", fontsize = 14)
```

```
[87]: Text(0.5, 1.0, 'Total Number of Titles with a G/TV-G Rating')
```



```
[88]: g_tv_g_amazon_only_df['rating'].value_counts()
```

```
[88]: rating
TV-G    282
G       245
Name: count, dtype: int64
```

```
[13]: plt.figure(figsize = (20.0, 10.0))
sns.histplot(finished_amazon_df, x='type')
```

```
[13]: <Axes: xlabel='type', ylabel='Count'>
```

